

Zhiyuan “Paul” Zhou

PERSONAL INFORMATION

Email	✉ zhiyuan.zhou@berkeley.edu
Website	🔗 https://zhouzypaul.github.io
Github	🐙 zhouzypaul
Linkedin	🌐 zhiyuan-paul-zhou
Google Scholar	📄 Zhiyuan Zhou

EDUCATION

UC Berkeley 2023 – Expected 2028

Ph.D. in Computer Science

- Researching robotics and deep reinforcement learning, advised by Prof. Sergey Levine
- College of Engineering Fellowship

Brown University 2019 – 2023

Sc.B. in Applied Mathematics and Computer Science

- GPA: 4.0 / 4.0; magna cum laude; CS Honors; Sigma Xi.
- Selected CS Coursework: Advanced Deep Learning, Machine Learning, Computer Vision, Collaborative Robotics, Multiprocessor Synchronization, Computer Systems.
- Selected Math Coursework: Computational Probability and Statistics, Pattern Theory, Statistics in Quantum Mechanics, Applied PDE & ODE, Honors Linear Algebra & Calculus.

PUBLICATIONS

**Efficient Online Reinforcement Learning Fine-Tuning
Need Not Retain Offline Data** 📄 2024

Zhiyuan Zhou*, Andy Peng*, Qiyang Li, Sergey Levine, Aviral Kumar
Preprint

**Autonomous Improvement of
Instruction Following Skills via Foundation Models** 📄 2024

Zhiyuan Zhou*, Pranav Atreya*, Abraham Lee, Homer Walke,
Oier Mees, Sergey Levine
Conference on Robot Learning (CoRL)

**Tiered Reward: Designing Rewards for
Specification and Fast Learning of Desired Behavior** 📄 2024

Zhiyuan Zhou, Shreyas Sundara Raman, Henry Sowerby, Michael L Littman
Reinforcement Learning Conference

Policy Transfer in Lifelong Reinforcement Learning through Learning Generalizing Features 📄 2023

Zhiyuan Zhou (Advisor: George Konidaris)
Undergraduate Honors Thesis

Characterizing the Action-Generalization Gap in Deep Q-Learning 📄 2022
Zhiyuan Zhou, Cameron Allen, Kavosh Asadi, George Konidakis
Multidisciplinary Conference on Reinforcement Learning and Decision Making (RLDM)

Designing Rewards for Fast Learning 📄 2022
Henry Sowerby, Zhiyuan Zhou, Michael L Littman
Multidisciplinary Conference on Reinforcement Learning and Decision Making (RLDM)
[Selected for oral]

ACADEMIC EXPERIENCE

RAIL, UC Berkeley 2023-present
Ph.D. student

- Working with Professor Sergey Levine on Deep Reinforcement Learning algorithms and finetuning robotic foundation models.

Intelligent Robot Lab, Brown University 2020-2023
Undergraduate Research Assistant

- Worked with professor George Konidakis and various Ph.D. students; one first-author paper accepted at RLDM and contributed to multiple ongoing papers.
- Researched various topics in deep Reinforcement Learning (RL): generalization and life-long learning; hierarchical RL through skill chaining; distributed hierarchical RL; action generalization in Deep RL.

RLAB, Brown University 2021-2023
Undergraduate Research Assistant

- Worked with professor Michael Littman; one first-author paper accepted at RLC, one at RLDM with oral, and one at ICML workshop.
- Researched various topics in reinforcement learning, and focused on the reward design for fast learning and behavior specification problem with formal guarantees.

Paradiso Lab, Brown University 2020
Undergraduate Research Assistant

- Worked with professor Michael Paradiso funded by the Brown Undergraduate Teaching and Research Award (UTRA); one first-author paper published in school research journal.
- Helped build a visual prosthetic device and researched topics in video object recognition.

INDUSTRY EXPERIENCE

Zencastr, Inc. July 2021 - Aug 2021
Natural Language Processing Engineer Intern

- Engineered and deployed a web app with websockets and FastAPI that allows users to edit (faulty) audio-to-text automatic transcriptions and provides a faster editing experience by intelligently recommending potentially incorrect segments: the recommendations are made by finding similar occurrences of user-made edits throughout the audio file with Keyword Spotting using language and acoustic models from Kaldi and Vosk-api.

- Sped up Keyword Spotting $2\times$ using multithreaded offline-decoding in Python and Shell; sped up automatic speech recognition $5\times$ using WeNet architecture (written in C++) and Speech Activity Detection models from Kaldi; model is pushed to production.
- Implemented a thread-safe MongoDB store with *asyncio* and *motor* to store user-made edits in the backend.

Zencastr, Inc.

Dec 2020 - Jan 2021

Machine Learning Engineer Intern

- Built a CNN in Keras that classifies audio files into speech, music, laughter, or noise with 93% accuracy; trained using audio data crawled from YouTube using *youtube-dl* and augmented by adding noise, changing pitch, and stretching time.
- Aligned audio-to-text transcriptions from DeepSpeech and Webspeech API using dynamic time warping and grapheme similarity.
- Built a Python library of machine learning utility scripts hosted privately on GitHub with Continuous Integration

TEACHING

Head Teaching Assistant

Spring 2022

CS1420 Machine Learning, Brown CS

- Managed a team of 20 teaching assistants and organized course logistics for 200 students.
- Built auto-grading pipelines for 12 coding assignments on Gradescope that enabled students to see code correctness shortly after handin.
- Answered questions through weekly TA hours and online discussion platform Edstem.

HONORS AND AWARDS

Berkeley College of Engineering Fellowship	2023
Sigma Xi Scientific Honors Society	2023
Brown University magna cum laude	2023
Brown CS Honors	2023
Brown CS Senior Prize	2023
UCSD ECE Fellowship (declined)	2023
Hack @ Brown Nelson Center for Entrepreneurship Award	2023
3 rd place in SELEF literary competition, STEM category	2022
Brown Undergraduate Teaching & Research Award	2021
Hack @ Brown Most Contrarian Hack & Wolfram Award	2021
Brown Undergraduate Teaching & Research Award	2019
227th (top 5%) in Putnam Math Competition, top 3 at Brown	2019
2nd Place in Hartshorn-Hypatia Math Contest	2019
Yongren Full Fellowship at PROMYS	2018
Provincial Top 1% in Chinese Physics Olympiad	2018
Regional Top 10 & International Top 100 in Physics Bowl	2018
Top 5% in AMC12	2018
Finalist in High School Mathematical Contest in Modeling	2017

INVITED TALKS

Designing Rewards for Fast Learning <i>Conference on Reinforcement Learning and Decision Making (RLDM)</i>	June 2022
Pareto Optimal Reward Functions <i>Robotics Lab, Brown CS</i>	July 2022